

TechNova Solutions | Evaluation of Github Copilot

Sharie Rhea

April 5, 2026

Document Overview

This document provides a high-level evaluation of the benefits, limitations, and risks of the agentic artificial intelligence (AI) tool Github Copilot to assess its feasibility for use in TechNova Solutions's development environment. The company's main requirements and operational constraints are detailed, followed by the report's overall recommendation. The evaluation section covers machine vs. deep learning, natural language processing, ethical risks, and topics including transparency, bias, intellectual property rights, security, and long-term impacts.

This evaluation was conducted through a review of academic literature, industry reports, and first-hand testing of Github Copilot and other AI tools. The final recommendation is based on a holistic view of the benefits and drawbacks of Copilot specifically for TechNova Solutions's use case.

Company Context

TechNova Solutions is a mid-sized software development company specializing in enterprise-grade solutions for clients in the healthcare and financial services industries. The company is exploring the integration of AI-powered code generation tools into its development workflow to enhance productivity, accelerate delivery, and maintain rigorous standards for code quality, security, and regulatory compliance.

Primary Services

- Custom enterprise software development for healthcare and financial institutions
- Development and maintenance of secure application programmer interfaces (APIs) for system integration and data exchange
- Data analytics and reporting platforms for regulated industries
- Ongoing software support, modernization, and cloud migration services
- Implementation of secure, compliant architectures aligned with industry regulations

Scale of Operations

- Serves dozens of enterprise clients across healthcare providers, insurers, and financial institutions
- Supports multiple large-scale applications with thousands of daily active users per client system
- Manages complex codebases across distributed development teams
- Handles sensitive data, including protected health information (PHI) and financial records
- Increasingly adopting DevOps and cloud-native practices to support scalability and rapid deployment

Compliance Requirements

- Health Insurance Portability and Accountability Act (HIPAA) – protection of patient health information
- Gramm-Leach-Bliley Act (GLBA) – safeguarding consumer financial data

- Payment Card Industry Data Security Standard (PCI DSS) – secure handling of payment card data
- Secure software development lifecycle (SDLC) practices, including code audits and vulnerability management
- Internal and client-driven compliance audits requiring traceability and documentation of code changes

Operational Constraints

- Must maintain high code quality and security standards despite increasing development velocity
- Needs transparency and auditability in development processes
- Requires scalable infrastructure to support continuous integration/continuous deployment (CI/CD) pipelines

Recommendation

The final recommendation of this report is: *GitHub Copilot is a worthwhile tool, but one that requires careful integration and planning.*

AI-powered code generation tools offer substantial potential to improve productivity and streamline development processes at TechNova Solutions. However, immediate, unrestricted adoption would introduce significant risks, particularly given the company's focus on healthcare and financial systems. Instead, a controlled and strategic implementation is strongly recommended. AI should be integrated incrementally, beginning with low-risk use cases such as documentation generation, boilerplate code creation, and debugging of isolated components.

To ensure safe and effective adoption, TechNova Solutions should implement a framework centered on human oversight and validation. All AI-generated code must undergo rigorous review, testing, and security analysis before being incorporated into production systems. Developers should be trained in prompt engineering techniques to improve output quality and reduce hallucinations, while also maintaining a critical mindset to avoid overreliance on AI-generated suggestions. Incorporating automated tools such as static analysis, security scanners, and license compliance checks will further mitigate risks.

AI should be viewed as an assistive tool rather than a replacement for human expertise. When implemented thoughtfully, it can enhance efficiency and innovation while maintaining the high standards required in regulated industries. By adopting a cautious, structured approach, TechNova Solutions can leverage AI to gain a competitive advantage while safeguarding its systems, clients, and reputation.

Evaluation

Machine Learning vs. Deep Learning

Machine learning (ML) is a branch of AI dedicated to developing algorithms that learn patterns from data to make predictions or decisions without explicit programming (Liu et al., 2025). ML is ideal for classification, anomaly detection, and regression tasks (Liu et al., 2025). Deep learning is a subset of ML inspired by the structure and function of the human brain that utilizes multi-layered neural networks with many hidden layers to automatically extract features from large and complex datasets (Liu et al., 2025).

From a performance standpoint, deep learning models offer greater flexibility and capabilities compared to traditional ML, especially in handling unstructured data such as documentation, user requirements, and code that is expressed in natural language. However, this increased performance comes at the cost of interpretability, as many deep learning models operate as “black boxes,” with little to no understanding of how specific outputs are generated (Liu et al., 2025).

Adopting Github Copilot to enhance developer efficiency, reduce coding errors, and streamline software development at TechNova Solutions involves accepting the tradeoff between the enhanced capabilities of deep learning neural networks and their lack of transparency.

Natural Language Processing

Modern AI code generation tools, such as Github Copilot, are based on large language models (LLMs). Unlike traditional ML systems, LLMs operate on natural language inputs using tokenization and vectorization. They excel in handling unstructured data and allow for a much more natural interaction with the AI tool. However, it is important to note that LLMs do not understand inputs, but rather generate outputs based on statistical likelihood (Liu et al., 2025).

This lack of understanding is the underlying cause for some of the main challenges of natural language processing models. LLMs, including Copilot, tend to produce authoritative responses even when there is uncertainty, leading to overconfidence (Liu et al., 2025). Additionally, AI may produce verbose responses if prompts are not carefully crafted, leading to increased cognitive load that decreases performance instead of boosting it.

The performance of LLMs is heavily influenced by prompt design. Well-structured prompts that include relevant context, constraints, and examples lead to a significant improvement in output quality (Liu et al., 2025). This means that engineers at TechNova Solutions must learn to effectively utilize AI tools from prompt to validation.

Comparison to Other AI Tools

Copilot is just one of many available AI tools, although it is tailored for code-centric tasks such as debugging, optimization, and step-by-step problem solving (Github, 2026). Copilot is an “agentic” tool, meaning that it is capable of performing actions within an environment other than a chat window. For example, Copilot works seamlessly with many popular integrated developer environments (IDEs) such as Visual Studio Code and Neovim to directly modify and test code. This makes Copilot well-suited to real-time coding assistance.

However, Copilot can struggle with higher-level tasks such as explaining code or brainstorming new ideas (Pandey et al., 2024). For these use cases, a tool such as OpenAI’s ChatGPT may be a better fit (OpenAI, 2026). No matter which tool is chosen, their effectiveness is highly dependent on prompt quality, iteration, and refinement along with hyperparameter tuning (Arora et al., 2024; Saini, 2025).

Real-World Problems in Various Domains

AI code generation tools are being used to address real-world problems across many domains. In healthcare, electronic health records systems, automatic data processing, and even diagnoses may be improved (R & C, 2025). In finance AI can optimize transaction processing, automate reporting and

auditing, enhance fraud detection, and trade stocks. Autonomous or self-driving cars are also a result of AI .

However, the integration of AI into various domains is challenging due to data sensitivity. The confidentiality and integrity of protected health information (PHI), personally identifiable information (PII), and financial records must be ensured (R & C, 2025). In safety-critical systems, such as self-driving cars and medical diagnoses, hallucinations or invalid decision-making pose a major risk.

Despite these challenges, AI tools such as Github Copilot still provide substantial benefits when used appropriately. The key is to enhance developer efficiency and accelerate development cycles using LLMs, while implementing robust validation processes to ensure AI-generated outputs are thoroughly tested and reviewed before deployment in real-world systems.

Attributes and Limitations

Agentic coding tools offer significant benefits, such as speed, scalability, and the ability to generate code and documentation rapidly across a wide range of domains. Coupled with this, AI tools are able to instantaneously analyze code, identify issues, and suggest improvements. This can lead to a significant reduction in the amount of time required for routine development tasks, such as writing documentation, creating boilerplate code, or researching APIs (Petryshyn & Lukoševičius, 2024).

A severe limitation of AI is the high likelihood of generating hallucinations. Syntactical, runtime execution, functional correctness, and quality cover the four categories of hallucinations in the context of code generation (Pavel et al., 2025). Syntactical and runtime execution issues are typically visible and easy to detect using automated systems or unit tests. Functional correctness and quality raise more subtle concerns that add to accumulating technical debt (Pavel et al., 2025). Hallucinations greatly increase the risk of bloated, poorly structured, or difficult to interpret code (Pavel et al., 2025). AI-generated code that makes it to production without proper review is likely to introduce security and optimization flaws (Faker & Toker, 2026).

Additionally, while AI tools perform well on small, self-contained tasks, they struggle with large, multi-file projects. There is a limit to the tokens and context that Copilot can handle, while a human developer is able to model an entire codebase and intelligently understand each moving part. This severely limits Copilot and other LLMs in enterprise environments with large legacy codebases.

Ethical Risks

AI poses significant ethical risks on the individual, societal, and environmental levels.

Unregulated AI can contribute to privacy erosion, job displacement, and mistrust (Liu et al., 2025; Yugendran et al., 2025). AI tools can introduce compounding issues in the form of ethical debt, occurring when “normative concerns are deferred or overlooked in favor of speed, functionality, or innovation” (Salah, 2025). For example, blind acceptance of code suggestions without proper privacy or security assessments and use of outdated or insecure practices due to time constraints constitute ethical debt. To compound the issue further, many programmers view AI tools as authoritative coauthors, leading to the erosion of responsibility and dilution of authorship (Salah, 2025).

Proper management, review, and oversight can reduce the amount of accrued ethical debt. Toolchain integrations that automatically review code for security, liability, and privacy violations can make these practices more feasible for resource-constrained groups (Salah, 2025). Users should also maintain a local

version of their code and data—stored with proper encryption—that is separate from versions used with third-party tools to maintain sensitive data. Personally identifiable information, trade secrets, and credentials should all be scrubbed before being exposed to an AI assistant (Sheggam & Zhang, 2024).

Impact of Bias

Ethical AI remains deeply challenging in practice, as models rely on formal rationality (Nishant et al., 2023). Humans, while inherently biased themselves, have the unique capability to shift between ethical theories. In this way, humans weigh competing values and adapt decisions to the individual context of every situation. Patterns, biases, or rules that work in one scenario do not necessarily apply in another (Nishant et al., 2023).

Hallucinations may inadvertently reinforce preheld misconceptions or biases (Virvou et al., 2025). AI tends to produce outputs with learned biases from the provided training data (Liu et al., 2025). On one hand, researchers are making great strides towards developing techniques to reduce or eliminate biases through algorithms and data analytics. However, there is a risk of overcorrection on the other. Efforts to “overclean” AI systems by eliminating all bias may unintentionally introduce opposing distortions. Humans naturally rely on pattern recognition shaped by repeated experiences, and attempting to remove every trace of bias may create new imbalances instead (Nishant et al., 2023).

Privacy and Transparency Risks

Interpretability is a major concern when using LLMs. When AI is used for decision making it is vital to understand how the decision was reached. In other words, the model should be transparent such that its output and validity may be properly assessed (Liu et al., 2025). Companies and individual developers should be sure to prioritize virtues such as transparency, fairness, accountability, ownership, dignity, and security in both policies and prompts (Huang et al., 2022). Additionally, prompts should be designed to produce interpretable results where the basis is understood (Liu et al., 2025). When possible, anonymize user data to protect privacy, or do not use AI tools on sensitive information. Implement safeguards to flag or prevent harmful output along with post-processing tools to enforce code quality standards. AI-generated code should be tagged or marked for proper accountability (Liu et al., 2025).

Privacy is another factor that may hinder the adoption of AI. Users are rightly concerned that the privacy of their data is not maintained, leading to information leakage in AI systems and potential violations of privacy regulations. Some industries, such as surveillance and e-marketing, are alienating potential consumers due to excessive hyperpersonalization, facial recognition, and other forms of privacy erosion (R & C, 2025; Yugendran et al., 2025). For this reason, responsible companies must clearly state how users’ data is stored and processed, including its interaction with AI tools.

Legal and Security Risks

Using AI may introduce significant intellectual property (IP) and security concerns. Violation of copyright—IP rights that an author intrinsically owns after the creation of a creative work—and licenses—rulesets that an author provides to utilize their IP in derivative works—pose a major risk for users of AI tools. AI models use IP in their training data and may also output content that is subject to IP protections (Lucchi, 2023). It is common for models to “memorize” aspects of open-source code libraries from their training data and regurgitate them, often with no indication that said information is

protected by a permissive or copyleft license (Herbold et al., 2024). In fact, AI is typically *incapable* of providing accurate information about the license status of generated code, meaning that simply telling the model to include the appropriate license is not enough (Xu et al., 2025). Copyleft licenses such as the GNU General Public License (GPL) present a challenge for corporate business models, as derivative works must also be distributed under the same terms (Black, 2023; Herbold et al., 2024). Strategies for mitigating IP concerns include static library or duplication detection, proper tagging of AI-generated code, and independent human review (Black, 2023; Github, 2026).

The question of whether AI-generated content is copyrightable is still being discussed. However, the current consensus is that content can hold IP rights when a human applies significant prompting, editing, or modification to constitute a creative work (Herbold et al., 2024; Lucchi, 2023).

Besides IP risks, security presents a challenge when integrating AI tools such as Github Copilot. LLMs often regurgitate data from training sets that include unsecure coding practices, hardcoded secrets or credentials, and poor encryption (Sheggam & Zhang, 2024). Without proper review, these vulnerabilities propagate to production codebases and cause major damage to customers and companies alike. Robust code validation, unit testing, integration testing, and—most importantly—human review, are vital for maintaining the high security standards required for the financial and healthcare sectors served by TechNova Solutions.

Long-Term Impact

Overreliance on AI tools may lead to skill degradation among developers, as critical thinking and problem-solving skills are replaced by automated suggestions. The accumulation of both technical and ethical debt in poorly understood codebases could lead to long-term maintenance challenges. This issue is particularly troublesome if AI-generated code is not properly tagged.

Another aspect to consider is evolving industry standards. As governments and organizations develop new guidelines, frameworks, and laws for AI usage, TechNova Solutions must be prepared to adapt to changing compliance requirements. This may involve ensuring transparency, accountability, and fairness along with data privacy and intellectual property protection.

Ultimately, the long-term success of integrating Copilot into TechNova Solutions depends on adopting a strategic approach to emphasize human oversight, ethical considerations, and adaptability to harness the benefits of AI responsibly.

References

- Arora, C., Sayeed, A. I., Licorish, S., Wang, F., & Treude, C. (2024). Optimizing Large Language Model Hyperparameters for Code Generation. *ArXiv (Cornell University)*.
<https://doi.org/10.48550/arxiv.2408.10577>
- Black, J. (2023, October 6). *5 Ways to Reduce GitHub Copilot Security and Legal Risks*. Fossa.
<https://fossa.com/blog/5-ways-to-reduce-github-copilot-security-and-legal-risks/>
- Faker, O., & Toker, S. (2026). Informed Skepticism: Navigating The Code Hallucination Dilemma. *2026 5th International Informatics and Software Engineering Conference (IISEC)*, Page 634-Page 639.
<https://doi.org/10.1109/iisec69317.2026.11418436>
- GitHub. (2026). *Copilot*. <https://github.com/features/copilot/ai-code-editor>
- Herbold, S., Valerius, B., Mojica-Hanke, A., Lex, I., & Mittel, J. (2024). Legal Aspects for Software Developers Interested in Generative AI Applications. *IEEE Software*, 1–7.
<https://doi.org/10.1109/ms.2024.3476677>
- Huang, C., Zhang, Z., Mao, B., & Yao, X. (2022). An Overview of Artificial Intelligence Ethics. *IEEE Transactions on Artificial Intelligence*, 4(4), 1–21. <https://doi.org/10.1109/tai.2022.3194503>
- Liu, A. C.-C., Law, O. M. K., & Law, I. (2025). *Machine Learning, Data Science Foundations with Python*. zyBook. <https://learn.zybooks.com/zybook/CS-530-10427.202617-1>
- Lucchi, N. (2023). ChatGPT: A Case Study on Copyright Challenges for Generative Artificial Intelligence Systems. *European Journal of Risk Regulation*, 15(3), 1–23. <https://doi.org/10.1017/err.2023.59>
- Nishant, R., Schneckenberg, D., & Ravishankar, M. N. (2023). The formal rationality of artificial intelligence-based algorithms and the problem of bias. *Journal of Information Technology*, 39(1).
<https://doi.org/10.1177/02683962231176842>
- OpenAI. (2026). *ChatGPT (GPT-5.2)*. <https://chat.openai.com/chat>
- Pandey, R., Singh, P., Wei, R., & Shankar, S. (2024). Transforming Software Development: Evaluating the Efficiency and Challenges of GitHub Copilot in Real-World Projects. *ArXiv (Cornell University)*.
<https://doi.org/10.48550/arxiv.2406.17910>
- Pavel, M., Petrovic, N., Mazur, L., Zolfaghari, V., Pan, F., & Knoll, A. (2025). Hallucination in LLM-Based Code Generation: An Automotive Case Study. *2025 3rd International Conference on Foundation and Large Language Models (FLLM)*, 101–106.
<https://doi.org/10.1109/fllm67465.2025.11391125>
- Petryshyn, B., & Lukoševičius, M. (2024). Optimizing Large Language Models for OpenAPI Code Completion. *ArXiv (Cornell University)*. <https://doi.org/10.48550/arxiv.2405.15729>
- R, L., & C, S. (2025). Social and Ethical Challenges of Artificial Intelligence in Surveillance and Security Systems. *2025 9th International Conference on Computational System and Information Technology for Sustainable Solutions (CSITSS)*, 1–6.
<https://doi.org/10.1109/csitss67709.2025.11294538>
- Saini, D. (2025, September 22). *Hyperparameters 101: Your One-Stop Guide for All Deep Learning Models (Part 1)*. Medium; Artificial Intelligence in Plain English.
<https://ai.plainenglish.io/hyperparameters-101-your-one-stop-guide-for-all-deep-learning-models-part-1-86f163f58a53>

- Salah, D. (2025). Invisible Risks, Visible Code: A Vision for Understanding Ethical Debt in AI-Based Coding. *2025 ACM/IEEE International Symposium on Empirical Software Engineering and Measurement (ESEM)*, 442–446. <https://doi.org/10.1109/esem64174.2025.00074>
- Sheggam H., & Zhang, X. (2024). *Exploring Security Risks and Mitigation Strategies in AI Code Helpers*. 1–6. <https://doi.org/10.1109/lisat63094.2024.10807934>
- Virvou, M., Tsihrintzis, G. A., & Tsichrintzi, E.-A. (2025). Hallucinations in Generative AI: Epistemic Risks for Learners in Educational Applications. *2025 16th International Conference on Information, Intelligence, Systems & Applications (IISA)*, 1–8. <https://doi.org/10.1109/iisa66859.2025.11311276>
- Xu, W., Gao, K., He, H., & Zhou, M. (2025). Licoeval: Evaluating LLMs on License Compliance in Code Generation. *ArXiv (Cornell University)*, 1665–1677. <https://doi.org/10.1109/icse55347.2025.00052>
- Yugendran V., Jenefa, L., Jayakumar, M., & M. Sakthivel. (2025). Ethics of AI Technology Interruption and Discomfort on Employee Turnover. 1–6. <https://doi.org/10.1109/icdsaai65575.2025.11011775>